

# Untitled

September 9, 2024

```
[337]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
pd.options.display.float_format='{:, .2f}'.format # to round all number to 0.2
sns.set()
```

About Dataset

Column Name

Description

sex

This column represents the gender of the individuals (female-male).

age

This column represents the age of the individuals in the dataset. Age is a crucial factor in assessing the risk of coronary heart disease.

education

This column represents the level of education of the individuals. It could be coded using categorical values indicating different levels of education attainment.

smokingStatus

This column likely represents the smoking status of the individuals, indicating whether they are smokers (yes) or non-smokers (no).

cigsPerDay

If an individual is a smoker, this column represents the number of cigarettes smoked per day.

BPMeds

This column indicates whether the individual is taking blood pressure medications (binary: 0 for not taking, 1 for taking).

prevalentStroke

This column indicates whether an individual has had a stroke prior to the study (binary: 0 for no, 1 for yes).

prevalentHyp

This column indicates whether an individual has hypertension (binary: 0 for no, 1 for yes).

diabetes

This column indicates whether an individual has diabetes (binary: 0 for no, 1 for yes).

totChol

This column represents the total cholesterol level of the individuals.

sysBP

This column represents the systolic blood pressure of the individuals.

diaBP

This column represents the diastolic blood pressure of the individuals.

BMI

This column represents the Body Mass Index (BMI) of the individuals, which is a measure of body fat based on height and weight.

heartRate

This column represents the resting heart rate of the individuals.

glucose

This column represents the fasting blood glucose level of the individuals.

CHDRisk

This column likely represents the Ten-Year Coronary Heart Disease (CHD) Risk for each individual, which is the target variable that you may want to predict or analyze.

```
[359]: data=pd.read_csv('Heart_Disease.csv')
data.head()
```

```
[359]:
```

	sex	age	education	smokingStatus	cigsPerDay	BPMeds	prevalentStroke	\
0	male	39	4	no	0	0	0	
1	female	46	2	no	0	0	0	
2	male	48	1	yes	20	0	0	
3	female	61	3	yes	30	0	0	
4	female	46	3	yes	23	0	0	

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	\
0		0	no	195	106.00	70.00	26.97	80	77
1		0	no	250	121.00	81.00	28.73	95	76
2		0	no	245	127.50	80.00	25.34	75	70
3		1	no	225	150.00	95.00	28.58	65	103
4		0	no	285	130.00	84.00	23.10	85	85

	CHDRisk
0	no

```

1      no
2      no
3      yes
4      no

```

## 1 Data\_Size

```
[360]: print(f'shape of data is {df.shape[0]} , {df.shape[1]}')
```

```
shape of data is 3674 , 16
```

```
[361]: df=data.copy()
```

## 2 Data\_Types

```
[362]: df.dtypes
```

```
[362]: sex          object
age             int64
education       int64
smokingStatus   object
cigsPerDay      int64
BPMeds          int64
prevalentStroke int64
prevalentHyp    int64
diabetes        object
totChol         int64
sysBP          float64
diaBP          float64
BMI            float64
heartRate      int64
glucose        int64
CHDRisk        object
dtype: object
```

```
[363]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3674 entries, 0 to 3673
Data columns (total 16 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   sex             3663 non-null   object
 1   age             3674 non-null   int64
 2   education       3674 non-null   int64
 3   smokingStatus   3661 non-null   object
 4   cigsPerDay      3674 non-null   int64

```

```

5   BPMeds           3674 non-null   int64
6   prevalentStroke  3674 non-null   int64
7   prevalentHyp     3674 non-null   int64
8   diabetes         3674 non-null   object
9   totChol          3674 non-null   int64
10  sysBP            3674 non-null   float64
11  diaBP            3674 non-null   float64
12  BMI              3674 non-null   float64
13  heartRate        3674 non-null   int64
14  glucose          3674 non-null   int64
15  CHDRisk          3674 non-null   object
dtypes: float64(3), int64(9), object(4)
memory usage: 459.4+ KB

```

### 3 Unique Values

```

[364]: for col in df.columns:
        print(f' {col}: \n number of unique value for each column {df[col].
↪unique()} , \n unique values is {df[col].unique()}')
        print('='*100)

sex:
number of unique value for each column 2 ,
unique values is ['male' 'female' nan]
=====

age:
number of unique value for each column 39 ,
unique values is [39 46 48 61 43 63 45 52 50 41 38 42 44 47 35 60 36 59 54 37
53 49 65 51
62 40 56 67 57 66 64 55 58 68 34 33 32 70 69]
=====

education:
number of unique value for each column 4 ,
unique values is [4 2 1 3]
=====

smokingStatus:
number of unique value for each column 2 ,
unique values is ['no' 'yes' nan]
=====

cigsPerDay:
number of unique value for each column 33 ,
unique values is [ 0 20 30 23 15 10  5 35 43  1 40  3  9  2 12  4 18 60 25 45
8 13 11  7

```

```

6 38 50 29 17 16 19 70 14]
=====
=====
BPMeds:
number of unique value for each column 2 ,
unique values is [0 1]
=====
=====
prevalentStroke:
number of unique value for each column 2 ,
unique values is [0 1]
=====
=====
prevalentHyp:
number of unique value for each column 2 ,
unique values is [0 1]
=====
=====
diabetes:
number of unique value for each column 2 ,
unique values is ['no' 'yes']
=====
=====
totChol:
number of unique value for each column 241 ,
unique values is [195 250 245 225 285 228 205 313 260 254 247 294 332 221 232
291 190 234
215 270 272 295 226 209 214 178 233 180 243 237 311 208 252 261 179 267
216 240 266 255 185 220 235 212 223 300 302 175 189 258 202 183 274 170
210 197 326 188 256 244 193 239 296 269 275 268 265 173 273 290 278 264
282 257 241 288 200 213 303 246 150 187 286 154 279 293 259 219 230 320
312 165 159 174 242 301 167 308 325 229 236 224 253 464 248 171 186 227
249 176 196 310 164 135 238 207 342 287 182 352 284 203 262 155 323 206
283 319 194 340 328 222 368 218 276 339 231 198 201 277 304 177 199 292
305 152 161 168 181 251 271 217 370 439 145 263 330 157 398 162 314 166
160 281 289 355 307 156 329 143 211 298 334 192 184 204 280 191 163 318
353 360 335 158 346 169 140 324 600 315 392 322 306 309 149 137 172 317
358 345 391 410 297 338 148 372 366 333 327 344 144 390 321 405 359 350
336 380 299 124 371 113 354 382 364 341 133 367 153 432 351 337 363 331
316 361 453 347 373 385 119]
=====
=====
sysBP:
number of unique value for each column 231 ,
unique values is [106. 121. 127.5 150. 130. 180. 138. 100. 141.5 162.
133. 131.
142. 124. 140. 112. 122. 139. 108. 148. 132. 137.5 102. 182.
115. 147. 124.5 160. 153. 111. 116.5 206. 96. 179.5 119. 116.

```

```

156.5 145. 114. 143.5 158. 157. 123.5 126.5 136. 154. 190. 107.
112.5 110. 138.5 155. 151. 152. 179. 113. 200. 132.5 126. 123.
134. 141. 135. 187. 127. 160.5 105. 109. 128. 118. 117.5 149.
180.5 136.5 212. 191. 121.5 173. 144. 129.5 117. 125. 144.5 170.
137. 94. 166. 177.5 129. 159. 130.5 107.5 189. 168. 197.5 146.
174. 98. 131.5 101. 158.5 97. 151.5 97.5 120. 204. 157.5 140.5
171. 215. 95. 156. 122.5 178. 146.5 113.5 197. 90. 109.5 165.
95.5 209. 162.5 295. 103. 134.5 115.5 174.5 163. 118.5 185. 220.
164. 120.5 98.5 161. 139.5 168.5 176. 163.5 128.5 167. 205.5 119.5
167.5 152.5 186. 183. 153.5 147.5 175. 142.5 192. 96.5 159.5 177.
102.5 244. 104. 213. 199. 184. 198. 114.5 125.5 111.5 105.5 143.
161.5 164.5 171.5 108.5 201. 148.5 172. 243. 145.5 187.5 99. 181.
133.5 100.5 135.5 172.5 103.5 149.5 182.5 186.5 217. 196. 193. 110.5
155.5 92. 169. 166.5 202. 150.5 195. 232. 85.5 184.5 188. 205.
169.5 210. 181.5 188.5 176.5 92.5 202.5 154.5 83.5 106.5 170.5 93.
175.5 207.5 199.5 101.5 248. 99.5 85. 230. 214. 192.5 104.5 194.
93.5 207. 185.5]

```

```

=====
=====

```

diaBP:

number of unique value for each column 142 ,

```

unique values is [ 70. 81. 80. 95. 84. 110. 71. 89. 107. 76.
88. 94.
90. 78. 84.5 70.5 82. 68. 91. 121. 85.5 85. 74. 92.5
98. 101. 73. 83.5 92. 63. 114. 77.5 69. 66. 82.5 102.
79. 75. 87. 99. 60. 67.5 72.5 106. 86.5 104. 86. 61.5
71.5 76.5 64. 77. 88.5 105. 96. 97. 100. 106.5 93. 80.5
124.5 61. 83. 67. 74.5 66.5 65. 72. 99.5 122.5 57. 57.5
111. 78.5 104.5 89.5 112. 55. 120. 118. 59. 133. 95.5 96.5
135. 64.5 68.5 98.5 62. 117. 59.5 103. 75.5 73.5 69.5 87.5
108. 93.5 90.5 114.5 62.5 94.5 140. 124. 91.5 115. 109. 102.5
65.5 105.5 103.5 63.5 79.5 107.5 142.5 109.5 58. 97.5 116.5 100.5
116. 119. 81.5 54. 132. 101.5 136. 51. 128. 125. 130. 110.5
113. 53. 108.5 112.5 52. 48. 56. 60.5 115.5 127.5]

```

```

=====
=====

```

BMI:

number of unique value for each column 1297 ,

unique values is [26.97 28.73 25.34 ... 26.7 43.67 19.71]

```

=====
=====

```

heartRate:

number of unique value for each column 72 ,

```

unique values is [ 80 95 75 65 85 77 60 79 76 93 72 98 64 70 71
62 73 90
96 68 63 88 78 83 100 84 57 50 74 86 55 92 66 87 110 81
56 89 82 54 69 67 52 61 140 130 58 104 94 105 91 53 108 106
59 107 48 112 125 103 44 47 45 97 122 102 120 99 115 143 101 46]

```

```

=====
=====
glucose:
number of unique value for each column 138 ,
unique values is [ 77 76 70 103 85 99 78 79 88 61 64 84 72 89 65
113 75 83
66 74 63 87 225 90 80 100 215 98 95 94 55 82 93 73 45 202
68 97 104 96 126 120 105 71 56 60 117 62 102 58 92 109 86 107
54 67 69 57 91 132 150 59 81 115 140 112 118 114 160 110 123 108
145 122 137 106 127 205 130 101 47 53 216 163 144 116 121 172 124 111
40 186 223 325 44 156 268 50 274 292 255 136 206 131 148 43 173 386
155 147 170 52 320 254 394 270 244 183 142 119 167 135 207 129 177 250
294 125 332 368 348 370 193 191 256 235 210 260]
=====
=====
CHDRisk:
number of unique value for each column 2 ,
unique values is ['no' 'yes']
=====
=====

```

Categorizing Age into Age Groups

```

[365]: df['age_group'] = pd.cut(df['age'], bins=[0, 30, 40, 60, 80], labels=['0-30', '30-40', '40-60', '60-80'])
df.drop(columns='age', inplace=True)

```

Converting Binary Medical Columns to Categorical Values for easier interpretation and analysis

```

[366]: df['BPMeds'] = df['BPMeds'].astype('object')
df['prevalentStroke'] = df['prevalentStroke'].astype('object')
df['prevalentHyp'] = df['prevalentHyp'].astype('object')

df['BPMeds'] = df['BPMeds'].replace([0, 1], ['no', 'yes'])
df['prevalentStroke'] = df['prevalentStroke'].replace([0, 1], ['no', 'yes'])
df['prevalentHyp'] = df['prevalentHyp'].replace([0, 1], ['no', 'yes'])

```

## 4 Show Data

```

[367]: df.head()

```

```

[367]:
   sex  education  smokingStatus  cigsPerDay  BPMeds  prevalentStroke \
0  male          4             no           0     no              no
1  female        2             no           0     no              no
2  male          1             yes          20     no              no
3  female        3             yes          30     no              no
4  female        3             yes          23     no              no

```

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	\
0	no	no	195	106.00	70.00	26.97	80	77	
1	no	no	250	121.00	81.00	28.73	95	76	
2	no	no	245	127.50	80.00	25.34	75	70	
3	yes	no	225	150.00	95.00	28.58	65	103	
4	no	no	285	130.00	84.00	23.10	85	85	

	CHDRisk	age_group
0	no	30-40
1	no	40-60
2	no	40-60
3	yes	60-80
4	no	40-60

```
[368]: df.tail()
```

```
[368]:
```

	sex	education	smokingStatus	cigsPerDay	BPMeds	prevalentStroke	\
3669	male	3	yes	25	no	no	
3670	male	3	yes	25	no	no	
3671	male	3	yes	25	no	no	
3672	male	3	yes	25	no	no	
3673	male	2	yes	25	no	no	

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	\
3669	no	no	208	137.50	82.50	25.58	75	63	
3670	no	no	208	137.50	82.50	25.58	75	63	
3671	no	no	208	137.50	82.50	25.58	75	63	
3672	no	no	208	137.50	82.50	25.58	75	63	
3673	no	no	208	137.50	82.50	25.97	69	68	

	CHDRisk	age_group
3669	yes	40-60
3670	yes	40-60
3671	yes	40-60
3672	yes	40-60
3673	yes	40-60

```
[369]: df.sample(2)
```

```
[369]:
```

	sex	education	smokingStatus	cigsPerDay	BPMeds	prevalentStroke	\
1573	male	1	yes	35	no	no	
1059	female	2	yes	20	no	no	

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	\
1573	no	no	188	120.00	82.50	31.67	80	68	
1059	no	no	149	122.00	72.00	21.30	85	75	



	CHDRisk	age_group
1573	no	40-60
1059	no	30-40

## 5 Missing\_\_Values

```
[370]: df.isna().sum()
```

```
[370]: sex                11
       education          0
       smokingStatus      13
       cigsPerDay          0
       BPMeds              0
       prevalentStroke     0
       prevalentHyp        0
       diabetes            0
       totChol             0
       sysBP               0
       diaBP               0
       BMI                 0
       heartRate           0
       glucose             0
       CHDRisk             0
       age_group           0
       dtype: int64
```

```
[371]: df=df.dropna()
```

```
[372]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 3652 entries, 0 to 3673
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sex                   3652 non-null  object
1   education             3652 non-null  int64
2   smokingStatus         3652 non-null  object
3   cigsPerDay            3652 non-null  int64
4   BPMeds                3652 non-null  object
5   prevalentStroke       3652 non-null  object
6   prevalentHyp          3652 non-null  object
7   diabetes              3652 non-null  object
8   totChol               3652 non-null  int64
9   sysBP                 3652 non-null  float64
10  diaBP                 3652 non-null  float64
11  BMI                   3652 non-null  float64
```

```

12 heartRate          3652 non-null    int64
13 glucose            3652 non-null    int64
14 CHDRisk            3652 non-null    object
15 age_group          3652 non-null    category
dtypes: category(1), float64(3), int64(5), object(7)
memory usage: 460.3+ KB

```

## 6 Duplicated\_Values

```

[374]: duplicated_data=df[df.duplicated(keep=False)]
       duplicated_data

```

```

[374]:      sex  education smokingStatus  cigsPerDay BPMeds prevalentStroke \
3118  male          3          yes          25      no          no
3658  male          3          yes          25      no          no
3659  male          3          yes          25      no          no
3660  male          3          yes          25      no          no
3661  male          3          yes          25      no          no
3662  male          3          yes          25      no          no
3663  male          3          yes          25      no          no
3664  male          3          yes          25      no          no
3665  male          3          yes          25      no          no
3666  male          3          yes          25      no          no
3667  male          3          yes          25      no          no
3668  male          3          yes          25      no          no
3669  male          3          yes          25      no          no
3670  male          3          yes          25      no          no
3671  male          3          yes          25      no          no
3672  male          3          yes          25      no          no

      prevalentHyp diabetes  totChol  sysBP  diaBP  BMI  heartRate  glucose \
3118          no      no      208  137.50  82.50  25.58          75          63
3658          no      no      208  137.50  82.50  25.58          75          63
3659          no      no      208  137.50  82.50  25.58          75          63
3660          no      no      208  137.50  82.50  25.58          75          63
3661          no      no      208  137.50  82.50  25.58          75          63
3662          no      no      208  137.50  82.50  25.58          75          63
3663          no      no      208  137.50  82.50  25.58          75          63
3664          no      no      208  137.50  82.50  25.58          75          63
3665          no      no      208  137.50  82.50  25.58          75          63
3666          no      no      208  137.50  82.50  25.58          75          63
3667          no      no      208  137.50  82.50  25.58          75          63
3668          no      no      208  137.50  82.50  25.58          75          63
3669          no      no      208  137.50  82.50  25.58          75          63
3670          no      no      208  137.50  82.50  25.58          75          63
3671          no      no      208  137.50  82.50  25.58          75          63
3672          no      no      208  137.50  82.50  25.58          75          63

```

	CHDRisk	age_group
3118	yes	40-60
3658	yes	40-60
3659	yes	40-60
3660	yes	40-60
3661	yes	40-60
3662	yes	40-60
3663	yes	40-60
3664	yes	40-60
3665	yes	40-60
3666	yes	40-60
3667	yes	40-60
3668	yes	40-60
3669	yes	40-60
3670	yes	40-60
3671	yes	40-60
3672	yes	40-60

```
[379]: df.duplicated().sum()
```

```
[379]: 0
```

```
[380]: df.drop_duplicates(inplace=True)
```

```
[381]: df.duplicated().any().sum()
```

```
[381]: 0
```

```
[382]: df.isna().sum()
```

```
[382]: sex          0
      education    0
      smokingStatus 0
      cigsPerDay    0
      BPMeds        0
      prevalentStroke 0
      prevalentHyp  0
      diabetes      0
      totChol       0
      sysBP         0
      diaBP         0
      BMI           0
      heartRate     0
      glucose       0
      CHDRisk       0
      age_group     0
```

```
dtype: int64
```

## 7 Statistical\_OverView

```
[384]: df.describe().T
```

```
[384]:
```

	count	mean	std	min	25%	50%	75%	max
education	3,637.00	1.98	1.02	1.00	1.00	2.00	3.00	4.00
cigsPerDay	3,637.00	9.03	11.91	0.00	0.00	0.00	20.00	70.00
totChol	3,637.00	236.88	44.13	113.00	206.00	234.00	263.00	600.00
sysBP	3,637.00	132.36	22.08	83.50	117.00	128.00	144.00	295.00
diaBP	3,637.00	82.90	11.96	48.00	75.00	82.00	90.00	142.50
BMI	3,637.00	25.79	4.06	15.54	23.08	25.38	28.04	56.80
heartRate	3,637.00	75.75	11.99	44.00	68.00	75.00	82.00	143.00
glucose	3,637.00	81.81	23.77	40.00	71.00	78.00	87.00	394.00

## 8 Analysis

```
[385]: df.columns
```

```
[385]: Index(['sex', 'education', 'smokingStatus', 'cigsPerDay', 'BPMeds',  
        'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP',  
        'diaBP', 'BMI', 'heartRate', 'glucose', 'CHDRisk', 'age_group'],  
        dtype='object')
```

1- univariate

Count Plots for Categorical Variables

```
[263]: plt.figure(figsize=(20,20))  
for i , col in enumerate(df.select_dtypes('object')):  
    plt.subplot(4,2,i+1)  
    ax=sns.countplot(x=df[col],order=df[col].value_counts().index)  
    ax.bar_label(ax.containers[0])# this line to show count for each col  
    plt.title(f'CountPlot of {col}')  
    plt.tight_layout()  
    print(f'Counts for {col}:\n', df[col].value_counts(), '\n')  
  
plt.show()
```

Counts for sex:

```
sex  
female    2026  
male      1611  
Name: count, dtype: int64
```

Counts for smokingStatus:

```
    smokingStatus
no      1857
yes     1780
Name: count, dtype: int64
```

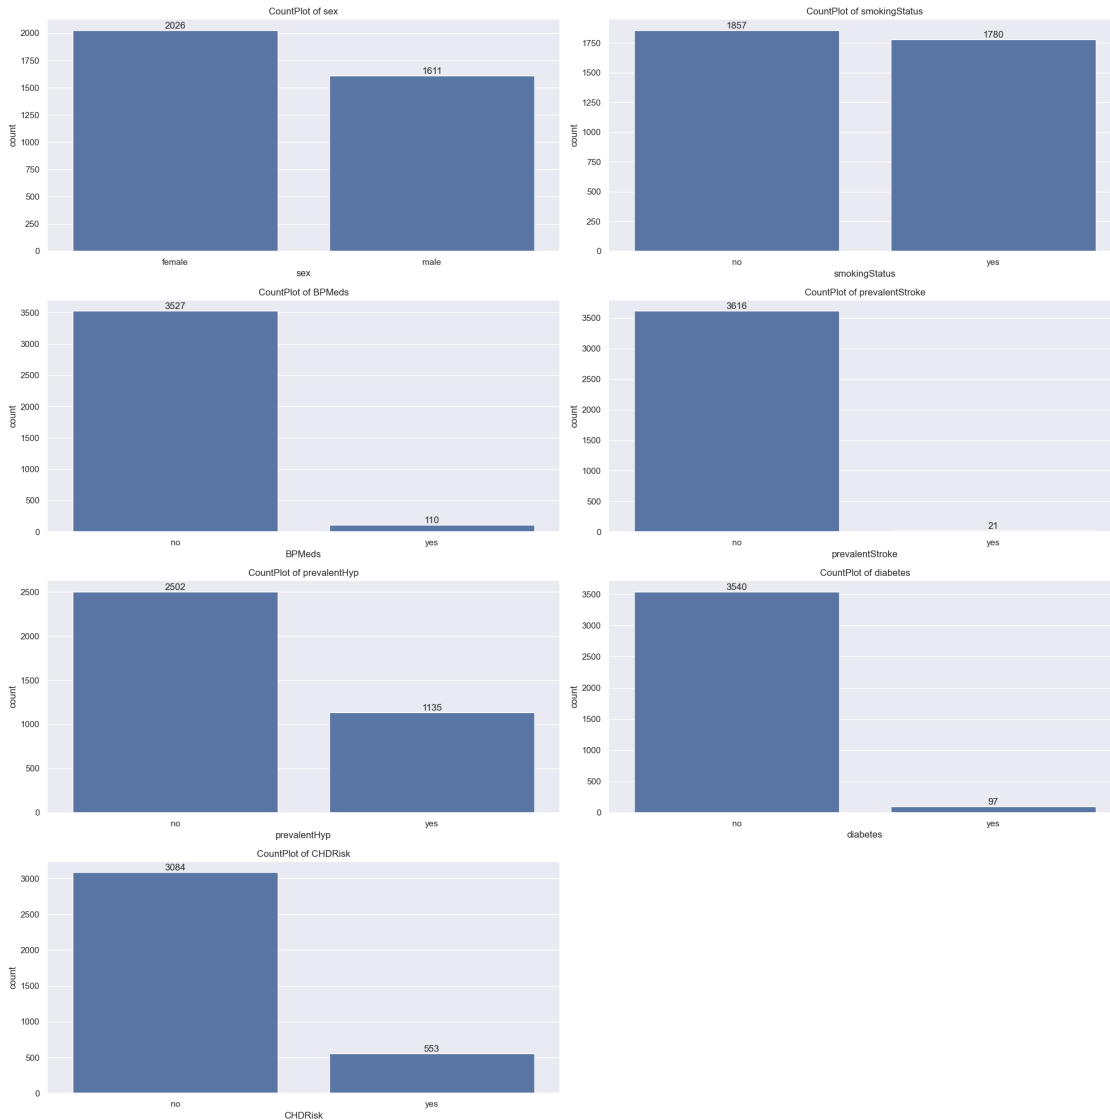
```
Counts for BPMeds:
    BPMeds
no      3527
yes      110
Name: count, dtype: int64
```

```
Counts for prevalentStroke:
    prevalentStroke
no      3616
yes       21
Name: count, dtype: int64
```

```
Counts for prevalentHyp:
    prevalentHyp
no      2502
yes     1135
Name: count, dtype: int64
```

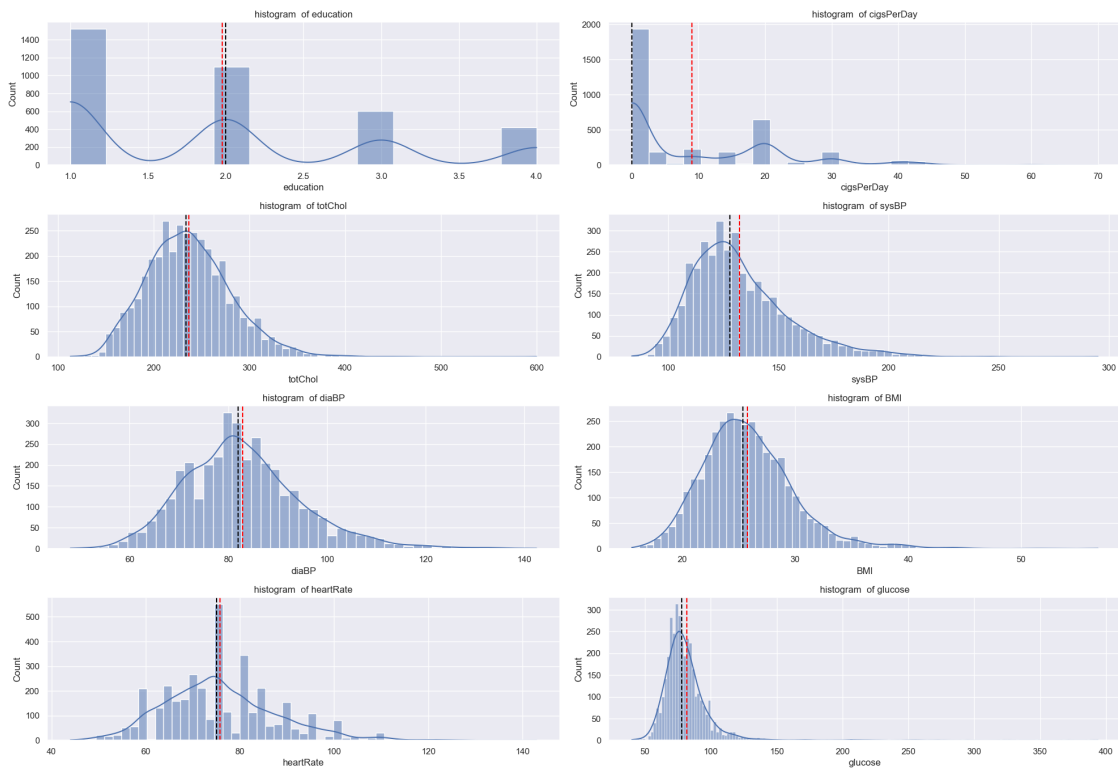
```
Counts for diabetes:
    diabetes
no      3540
yes       97
Name: count, dtype: int64
```

```
Counts for CHDRisk:
    CHDRisk
no      3084
yes      553
Name: count, dtype: int64
```




Box Plots for numerical columns

```
[386]: plt.figure(figsize=(20,20))
for i , col in enumerate(df.select_dtypes('number')):
    plt.subplot(6,2,i+1)
    ax=sns.histplot(df[col],kde=True)
    ax.axvline(df[col].mean(),color='red',linestyle='--')
    ax.axvline(df[col].median(),color='black',linestyle='--')
    plt.title(f'histogram of {col}')
    plt.tight_layout()
plt.show()
```




data is approximately normally distributed bec mean is closer to Median in histogram and descriptive stat

```
[387]: plt.figure(figsize=(20,40))
for i , col in enumerate(df.select_dtypes('number')):
    plt.subplot(2,5,i+1)
    ax=sns.boxplot(df[col])
    plt.title(f'boxplot of {col}')
    plt.tight_layout()
plt.show()
```





```
[388]: age_grouped_valueCount=df['age_group'].value_counts()
# age_grouped_valueCount
sns.barplot(x=age_grouped_valueCount.index,y=age_grouped_valueCount.
↪values,hue=age_grouped_valueCount)
plt.title('Distribution of Age Groups')
plt.xlabel('Age Groups')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.show()
```



## 2- Bivariate

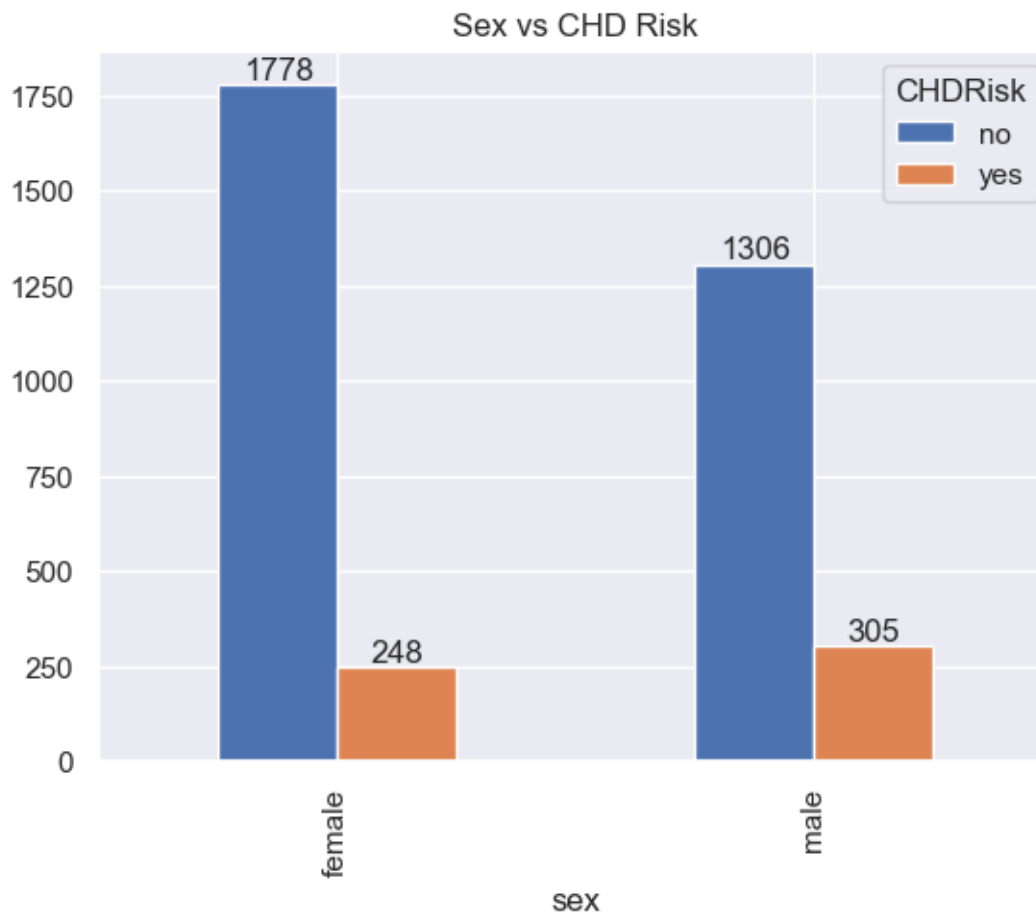
### a. Categorical vs Categorical:

```
[389]: crossTab_df=pd.crosstab(df['sex'],df['CHDRisk'])
print(crossTab_df)
```

```
ax=crossTab_df.plot(kind='bar',title='Sex vs CHD Risk')
ax.bar_label(ax.containers[0])
ax.bar_label(ax.containers[1])

plt.show()
```

CHDRisk	no	yes
sex		
female	1778	248
male	1306	305

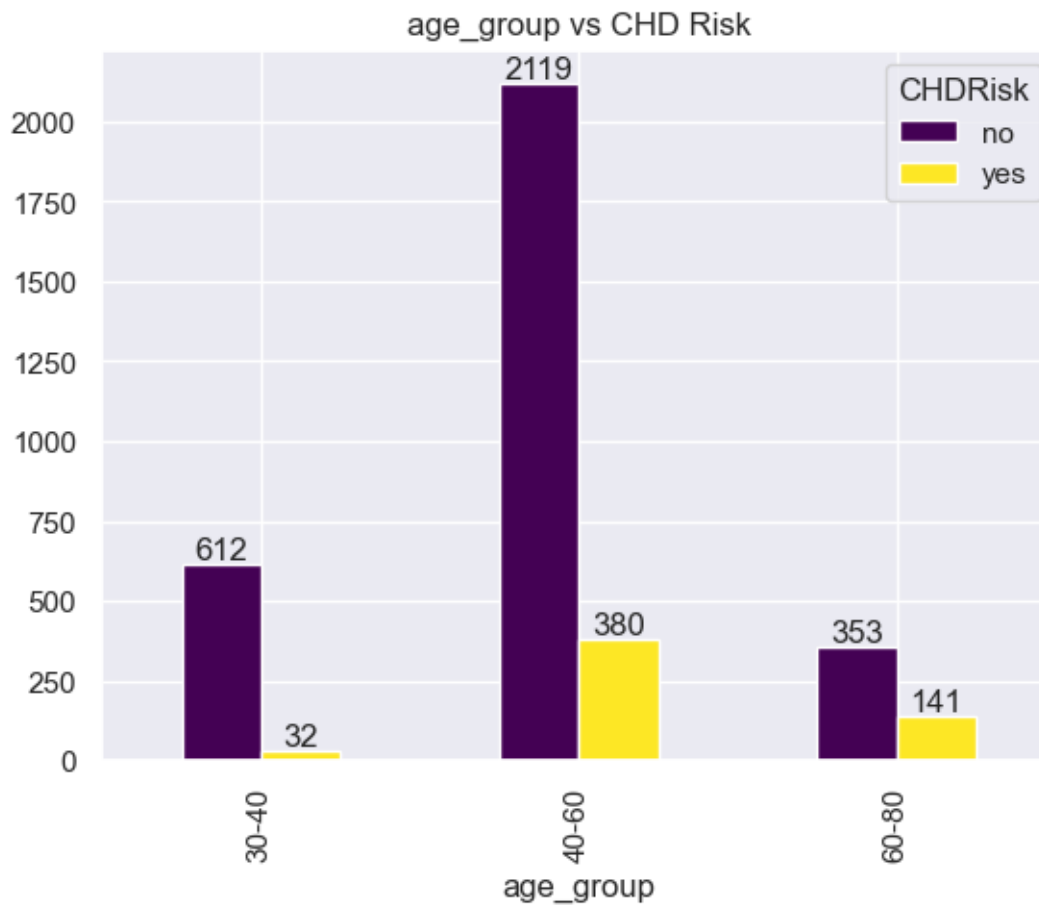


```
[390]: crossTab_df=pd.crosstab(df['age_group'],df['CHDRisk'])
print(crossTab_df)

ax=crossTab_df.plot(kind='bar',title='age_group vs CHD Risk',colormap='viridis')
ax.bar_label(ax.containers[0])
ax.bar_label(ax.containers[1])
```

```
plt.show()
```

CHDRisk	no	yes
age_group		
30-40	612	32
40-60	2119	380
60-80	353	141

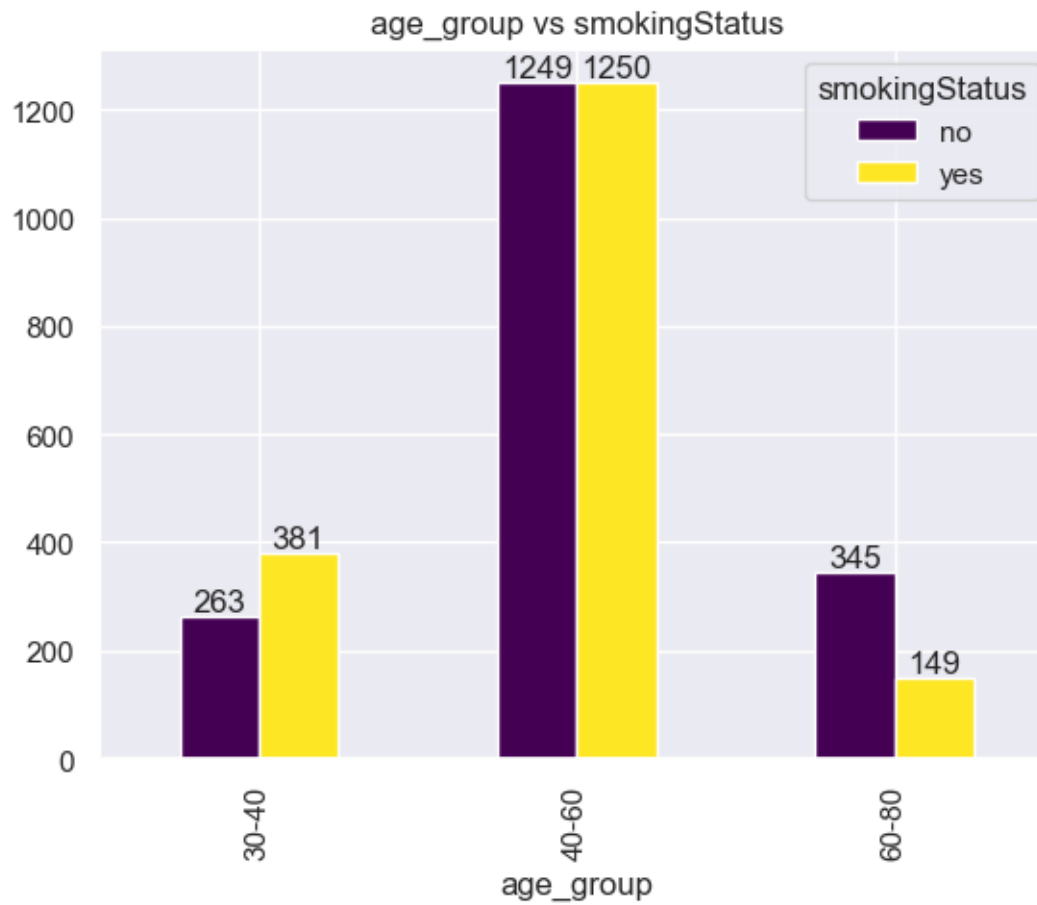


```
[391]: crossTab_df=pd.crosstab(df['age_group'],df['smokingStatus'])
print(crossTab_df)

ax=crossTab_df.plot(kind='bar',title='age_group vs_
↳smokingStatus',colormap='viridis')
ax.bar_label(ax.containers[0])
ax.bar_label(ax.containers[1])

plt.show()
```

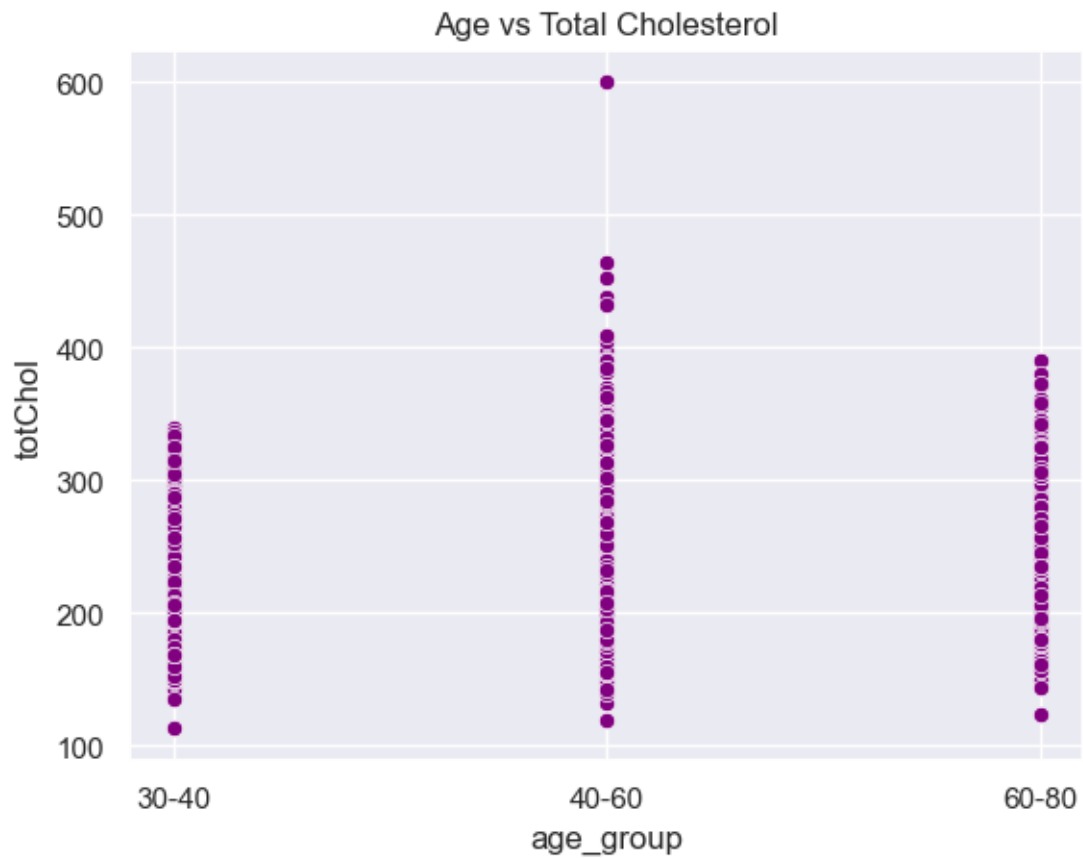
smokingStatus	no	yes
age_group		
30-40	263	381
40-60	1249	1250
60-80	345	149



b. Numerical vs Categorical

```
[392]: sns.scatterplot(data=df,x='age_group',y='totChol',color='purple')
plt.title('Age vs Total Cholesterol')
```

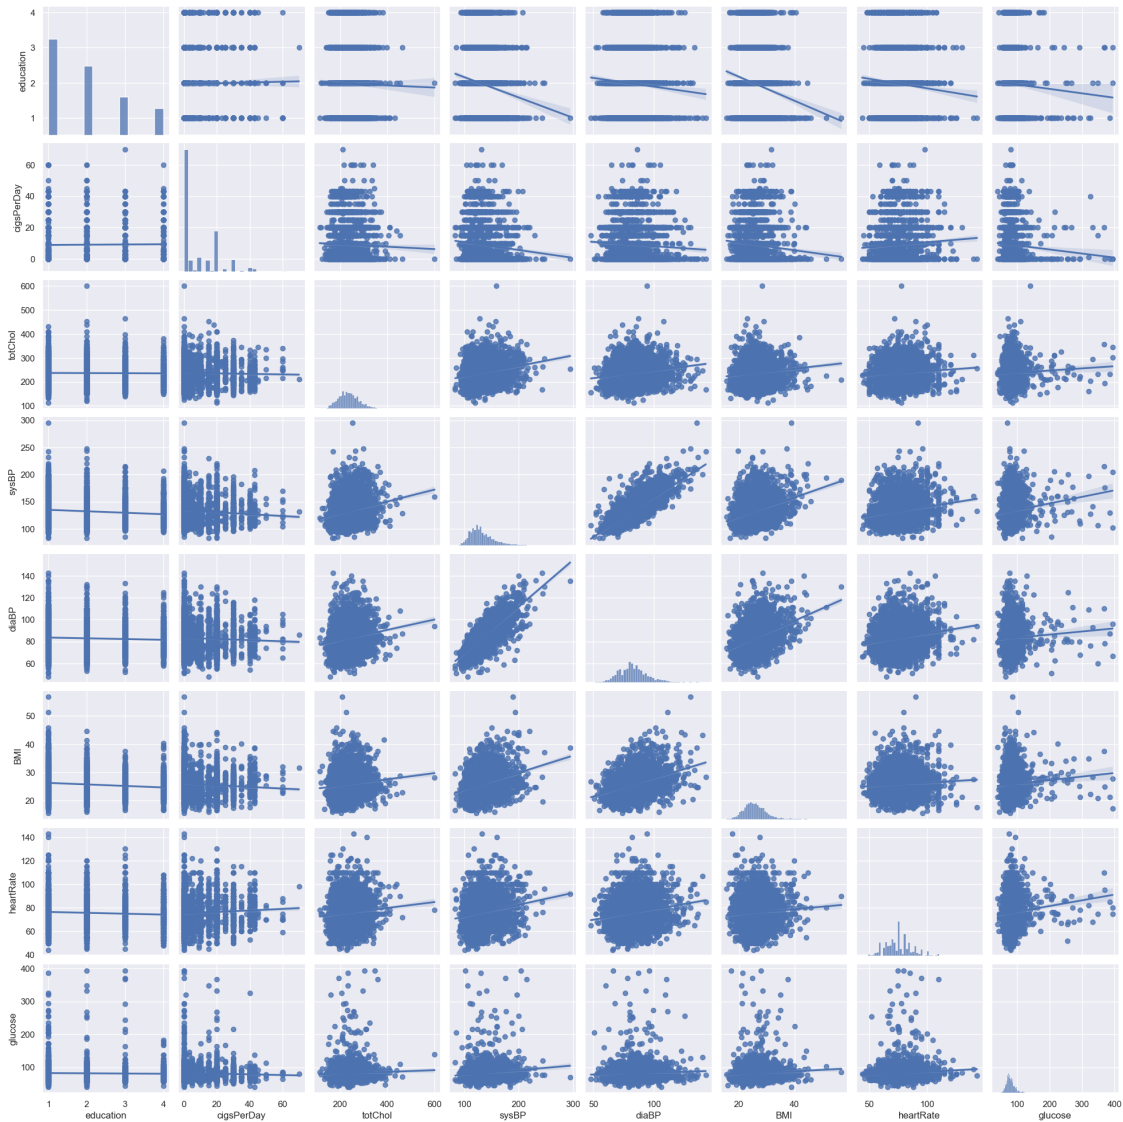
```
[392]: Text(0.5, 1.0, 'Age vs Total Cholesterol')
```



### 3- Multivariate Analysis

```
[393]: numerical_columns=df.select_dtypes('number')
sns.pairplot(data=numerical_columns,kind='reg')
```

```
[393]: <seaborn.axisgrid.PairGrid at 0x242ab284fb0>
```




```
[394]: # Generate a mask for the upper triangle
mask = np.zeros_like(numerical_columns.corr())
mask[np.triu_indices_from(mask)] = True

plt.figure(figsize=(10, 8))

sns.heatmap(numerical_columns.corr(),
            annot=True,
            vmin=0.1,
            mask=mask,
            linewidths=3.5,
            linecolor='white',
            cmap='coolwarm'),
```

*# Set the minimum correlation value color*  
*# Apply the mask to hide the upper triangle*  
*# Increase the space between squares*  
*# Set the line color between the squares*

```
plt.show()
```



there is a strong correlation between sysBP and diaBp

## 9 Health Metrics: Normal, Borderline, and High Ranges

Metric	Normal Range	Borderline/Elevated	High/Abnormal
<b>Total Cholesterol (totChol)</b>	Less than 200 mg/dL	200-239 mg/dL	240 mg/dL and above
<b>Systolic Blood Pressure (sysBP)</b>	90-120 mmHg	120-129 mmHg	130 mmHg and above

Metric	Normal Range	Borderline/Elevated High/Abnormal	
<b>Diastolic Blood Pressure (diaBP)</b>	60-80 mmHg	80-89 mmHg	90 mmHg and above
<b>Body Mass Index (BMI)</b>	18.5-24.9	25-29.9	30 and above (Obesity)
<b>Heart Rate (heartRate)</b>	60-100 bpm	-	Less than 60 bpm (Bradycardia), more than 100 bpm (Tachycardia)
<b>Glucose</b>	70-99 mg/dL (fasting)	100-125 mg/dL (Prediabetes)	126 mg/dL and above (Diabetes)

Number of Individuals Classified as Healthy Based on Key Health Metrics

```
[395]: healthy_conditions=(
        (df['sex'].isin(['male','female']))&
        (df['smokingStatus'].isin(['no','yes']))&
        (df['cigsPerDay'].between(0,20))&
        (df['BPMeds']=='no')&
        (df['prevalentStroke']=='no')&
        (df['prevalentHyp']=='no')&
        (df['diabetes']=='no')&
        (df['diaBP'].between(60,80))&
        (df['sysBP'].between(90,120))&
        (df['BMI'].between(18.5,24.9))&
        (df['heartRate'].between(60,100))&
        (df['glucose'].between(70,99))&
        (df['totChol']<200)
    )
    healthy_People=df[healthy_conditions]
    length_healthy_people = len(healthy_People)
    print(f"Number of healthy people: {length_healthy_people}")
```

Number of healthy people: 100

```
[398]: healthy_People.sample(10)
```

```
[398]:      sex  education  smokingStatus  cigsPerDay  BPMeds  prevalentStroke  \
530   female         2             no           0      no             no
1469  female         1             no           0      no             no
1729  female         4             no           0      no             no
635   female         1            yes          20      no             no
636   female         3             no           0      no             no
1452  female         1            yes          20      no             no
3524   male         4            yes          20      no             no
1778  female         1            yes          20      no             no
1460  female         2            yes          10      no             no
```



1381	female	3	yes	3	no	no
------	--------	---	-----	---	----	----

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	\
530	no	no	170	98.50	69.50	19.64	71	77	
1469	no	no	170	110.00	69.00	23.48	75	83	
1729	no	no	150	108.00	70.50	20.42	72	88	
635	no	no	168	117.00	74.00	21.51	67	77	
636	no	no	197	107.00	73.00	19.78	63	76	
1452	no	no	175	117.50	73.50	22.15	65	75	
3524	no	no	198	116.00	74.00	23.99	75	78	
1778	no	no	166	112.00	73.50	21.64	75	93	
1460	no	no	169	119.00	72.00	19.78	60	74	
1381	no	no	186	114.00	77.00	21.01	80	85	

	CHDRisk	age_group
530	no	30-40
1469	no	30-40
1729	no	30-40
635	no	30-40
636	no	40-60
1452	no	30-40
3524	no	40-60
1778	yes	30-40
1460	no	40-60
1381	no	30-40

index of 1778 write yes in CHDRisk Although it is among the healthy people

## 10 Distribution and Classification of Health Metrics and CHDRisk

- (barplot)

```
[399]: from matplotlib.ticker import FixedLocator

# Create dictionaries for each category
Smokers = {'NonSmokers': df[(df.cigsPerDay == 0)][ 'cigsPerDay'].count(),
          'Smokers': df[(df.cigsPerDay > 0) & (df.cigsPerDay <= 20)][ 'cigsPerDay'].count(),
          'Dangerous Smokers': df[(df.cigsPerDay > 20) & (df.cigsPerDay <= 50)][ 'cigsPerDay'].count(),
          'High Smokers': df[(df.cigsPerDay > 50)][ 'cigsPerDay'].count()}

Chol = {'LowChol': df[(df.totChol < 125)][ 'totChol'].count(),
        'NormalChol': df[(df.totChol > 125) & (df.totChol <= 200)][ 'totChol'].count(),
```

```

    'DangerousChol': df[(df.totChol > 200) & (df.totChol <= 239)]['totChol'].count(),
    'HighChol': df[(df.totChol > 240)]['totChol'].count()})

SysBP = {'LowsysBP': df[(df.sysBP < 120)]['sysBP'].count(),
        'NormalsysBP': df[(df.sysBP >= 120) & (df.sysBP <= 129)]['sysBP'].count(),
        'DangeroussysBP': df[(df.sysBP >= 130) & (df.sysBP <= 179)]['sysBP'].count(),
        'HighsysBP': df[(df.sysBP >= 180)]['sysBP'].count()})

DiaBP = {'LowdiaBP': df[(df.diaBP < 80)]['diaBP'].count(),
        'NormaldiaBP': df[(df.diaBP >= 80) & (df.diaBP <= 84)]['diaBP'].count(),
        'DangerousdiaBP': df[(df.diaBP >= 85) & (df.diaBP <= 109)]['diaBP'].count(),
        'HighdiaBP': df[(df.diaBP >= 110)]['diaBP'].count()})

BMI = {'LowBMI': df[(df.BMI < 18.5)]['BMI'].count(),
        'NormalBMI': df[(df.BMI >= 18.5) & (df.BMI <= 29.9)]['BMI'].count(),
        'DangerousBMI': df[(df.BMI >= 30) & (df.BMI <= 39.9)]['BMI'].count(),
        'HighBMI': df[(df.BMI >= 40)]['BMI'].count()})

HeartRate = {'LowHR': df[(df.heartRate < 60)]['heartRate'].count(),
        'NormalHR': df[(df.heartRate >= 60) & (df.heartRate <= 100)]['heartRate'].count(),
        'DangerousHR': df[(df.heartRate > 100)]['heartRate'].count()})

Glo = {'LowGol': df[(df.glucose < 70)]['glucose'].count(),
        'NormalGol': df[(df.glucose >= 70) & (df.glucose <= 100)]['glucose'].count(),
        'DangerousGol': df[(df.glucose > 100)]['glucose'].count()})

BPMeds = {'Yes': df[(df.BPMeds == 'yes')]['BPMeds'].count(),
        'No': df[(df.BPMeds == 'no')]['BPMeds'].count()})

prevalentStroke = {'Yes': df[(df.prevalentStroke == 'yes')]['prevalentStroke'].count(),
        'No': df[(df.prevalentStroke == 'no')]['prevalentStroke'].count()})

prevalentHyp = {'Yes': df[(df.prevalentHyp == 'yes')]['prevalentHyp'].count(),
        'No': df[(df.prevalentHyp == 'no')]['prevalentHyp'].count()})

gender = {'Male': df[(df['sex'] == 'male')]['sex'].count(),

```

```

        'Female': df[df['sex'] == 'female']['sex'].count()}

CHDRisk = {'Yes': df[(df.CHDRisk == 'yes')]['CHDRisk'].count(),
          'No': df[(df.CHDRisk == 'no')]['CHDRisk'].count()}

Diabetes = {'Yes': df[(df.diabetes == 'yes')]['diabetes'].count(),
          'No': df[(df.diabetes == 'no')]['diabetes'].count()}

# List of dictionaries and titles
dicts = [Smokers, Chol, SysBP, DiaBP, BMI, HeartRate, Glo, BPMeds,
        ↪prevalentStroke, prevalentHyp, Diabetes, gender, CHDRisk]
dicts_title = ['Smokers', 'Cholesterol', 'Systolic BP', 'Diastolic BP', 'BMI',
        ↪'Heart Rate', 'Glucose', 'BP Meds', 'Prevalent Stroke', 'Prevalent
        ↪Hypertension', 'Diabetes', 'Gender', 'CHD Risk']

# Plotting
plt.figure(figsize=(20, 20))

for e, dict_data in enumerate(dicts):
    d = pd.DataFrame(dict_data.items(), columns=['Classification', 'Value'])

    print(f" {dicts_title[e]}")
    print(d)
    print('=' * 50)

    plt.subplot(5, 3, e+1)
    ax = sns.barplot(data=d, x='Classification', y='Value')
    ax.bar_label(ax.containers[0])# to show count outside plot
    ax.set_title(dicts_title[e])
    ax.set_xlabel('')
    ax.set_ylabel('')

    ax.set_xticks(range(len(d['Classification'])))
    ax.set_xticklabels(d['Classification'], rotation=62, ha='right')

plt.tight_layout()
plt.show()

```

	Classification	Value
0	NonSmokers	1857
1	Smokers	1385
2	Dangerous Smokers	385
3	High Smokers	10

```

=====
Cholesterol
  Classification  Value
0      LowChol    3
1      NormalChol 755
2      DangerousChol 1241
3      HighChol   1570
=====

Systolic BP
  Classification  Value
0      LowsysBP   1100
1      NormalsysBP 798
2      DangeroussysBP 1585
3      HighsysBP   143
=====

Diastolic BP
  Classification  Value
0      LowdiaBP   1430
1      NormaldiaBP 726
2      DangerousdiaBP 1374
3      HighdiaBP   96
=====

BMI
  Classification  Value
0      LowBMI     49
1      NormalBMI  3123
2      DangerousBMI 429
3      HighBMI     21
=====

Heart Rate
  Classification  Value
0      LowHR      175
1      NormalHR   3375
2      DangerousHR 87
=====

Glucose
  Classification  Value
0      LowGol     712
1      NormalGol  2637
2      DangerousGol 288
=====

BP Meds
  Classification  Value
0      Yes        110
1      No         3527
=====

Prevalent Stroke
  Classification  Value

```

0	Yes	21
1	No	3616

=====

Prevalent Hypertension		
	Classification	Value
0	Yes	1135
1	No	2502

=====

Diabetes		
	Classification	Value
0	Yes	97
1	No	3540


=====

Gender		
	Classification	Value
0	Male	1611
1	Female	2026

=====


CHD Risk		
	Classification	Value
0	Yes	553
1	No	3084

=====



- (pie)

```
[400]: plt.figure(figsize=(20,20))
for i , dict_col in enumerate(dict):
    d=pd.DataFrame(dict_col.items(),columns=['Classification','Value'])
    plt.subplot(6,3,i+1)
    plt.pie(d['Value'],labels=d['Classification'],autopct='%1.1f%%')
    plt.tight_layout()
    plt.title(f'\n \n {dicts_title[i]}')
plt.show()
```



```
[401]: non_healthy_people=len(df)-length_healthy_people
non_healthy_people
```


[401]: 3537

Distribution of Healthy vs. Non-Healthy Individuals

```
[402]: size=[length_healthy_people,non_healthy_people]
labels=[' \n Healthy','Non-Healthy']
colors = ['#66b3ff', '#ff9999']

plt.figure(figsize=(6,6))
plt.pie(size,labels=labels,colors=colors,autopct='%1.1f%%',startangle=90,
        explode=(0, 0.1))
plt.title('Distribution of Healthy and Non-Healthy People')
```

```
plt.axis('equal')
plt.show()
```



```
[403]: df.columns
```

```
[403]: Index(['sex', 'education', 'smokingStatus', 'cigsPerDay', 'BPMeds',
            'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP',
            'diaBP', 'BMI', 'heartRate', 'glucose', 'CHDRisk', 'age_group'],
            dtype='object')
```

percentage of males and females with various health conditions

```
[405]: disease_columns=['smokingStatus', 'BPMeds', 'prevalentStroke', 'prevalentHyp',
                        ↪ 'diabetes', 'CHDRisk']

plt.figure(figsize=(15,10))
```



```

for i ,col in enumerate(disease_columns):
    gender_disease_count=df.groupby(['sex',col]).size().
    ↪reset_index(name='count')
#     print(gender_disease_count)
    gender_total=df.groupby('sex').size().reset_index(name='total')
#     print(gender_total)
    merged_df=pd.merge(gender_disease_count,gender_total,on='sex')
#     print(merged_df)
    merged_df['Percentage']=(merged_df['count'] / merged_df['total'])*100
#     print(merged_df)
    print(f"\nPercentage of Males and Females with {col}:")

    #iterrows() function returns both the index and the row data.
    #If you are only interested in the row and do not need the index, you can
    ↪use _ to indicate that you are intentionally ignoring the index.
    for _, row in merged_df.iterrows():
        print(f"Gender: {row['sex']}, Disease Status: {row[col]}, Percentage:
        ↪{row['Percentage']:.2f}%")

    plt.subplot(3, 3, i+1)
    ax=sns.barplot(data=merged_df, x='sex', y='Percentage', hue=col)
    ax.bar_label(ax.containers[0])# to show count outside plot
    ax.bar_label(ax.containers[1])

    ax.set_title(f'Percentage of Males and Females with {col}')
    ax.set_ylabel('')
    ax.set_xlabel('')

plt.tight_layout()
plt.show()

```

Percentage of Males and Females with smokingStatus:

Gender: female, Disease Status: no, Percentage: 60.32%  
 Gender: female, Disease Status: yes, Percentage: 39.68%  
 Gender: male, Disease Status: no, Percentage: 39.42%  
 Gender: male, Disease Status: yes, Percentage: 60.58%

Percentage of Males and Females with BPMeds:

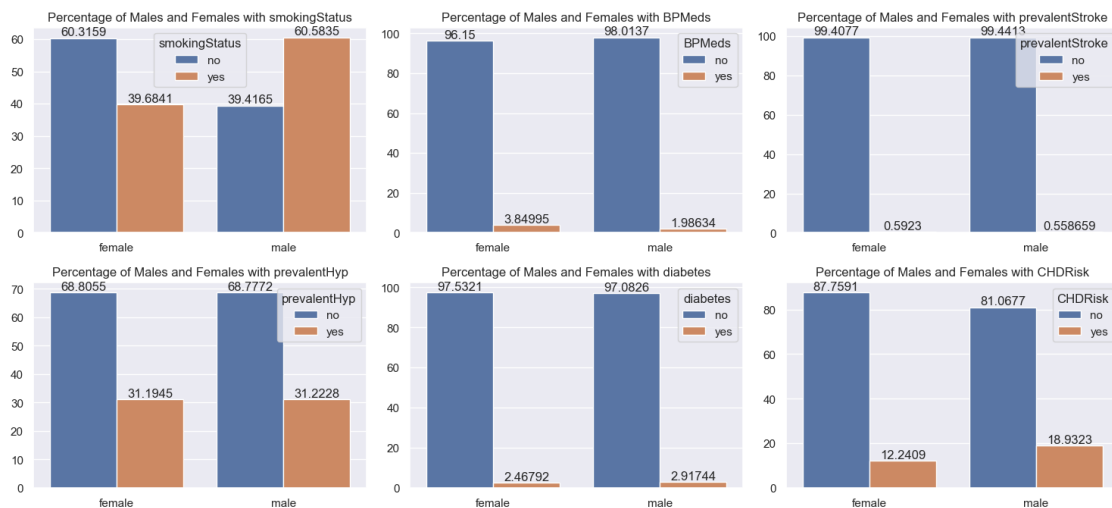
Gender: female, Disease Status: no, Percentage: 96.15%  
 Gender: female, Disease Status: yes, Percentage: 3.85%  
 Gender: male, Disease Status: no, Percentage: 98.01%  
 Gender: male, Disease Status: yes, Percentage: 1.99%

Percentage of Males and Females with prevalentStroke:  
 Gender: female, Disease Status: no, Percentage: 99.41%  
 Gender: female, Disease Status: yes, Percentage: 0.59%  
 Gender: male, Disease Status: no, Percentage: 99.44%  
 Gender: male, Disease Status: yes, Percentage: 0.56%

Percentage of Males and Females with prevalentHyp:  
 Gender: female, Disease Status: no, Percentage: 68.81%  
 Gender: female, Disease Status: yes, Percentage: 31.19%  
 Gender: male, Disease Status: no, Percentage: 68.78%  
 Gender: male, Disease Status: yes, Percentage: 31.22%

Percentage of Males and Females with diabetes:  
 Gender: female, Disease Status: no, Percentage: 97.53%  
 Gender: female, Disease Status: yes, Percentage: 2.47%  
 Gender: male, Disease Status: no, Percentage: 97.08%  
 Gender: male, Disease Status: yes, Percentage: 2.92%

Percentage of Males and Females with CHDRisk:  
 Gender: female, Disease Status: no, Percentage: 87.76%  
 Gender: female, Disease Status: yes, Percentage: 12.24%  
 Gender: male, Disease Status: no, Percentage: 81.07%  
 Gender: male, Disease Status: yes, Percentage: 18.93%




relationship between various health conditions and CHDRisk among females

```
[406]: disease = df[['sex', 'smokingStatus', 'BPMeds', 'prevalentStroke',
                    ↪ 'prevalentHyp', 'diabetes', 'CHDRisk']]

gender_female=disease[disease.sex == 'female']
```

```
plt.figure(figsize=(20,20))
for i ,col in enumerate(gender_female):
    plt.subplot(4,3,i+1)
    ax=sns.countplot(data=gender_female,x=col ,hue='CHDRisk')
    ax.bar_label(ax.containers[0])
    ax.bar_label(ax.containers[1])
    ax.set_title('\n\n' + col)
    ax.set_xlabel('')
    ax.set_ylabel('')
```




relationship between various health conditions and CHDRisk among males

```
[407]: disease = df[['sex','smokingStatus', 'BPMeds', 'prevalentStroke',
    ↪ 'prevalentHyp', 'diabetes', 'CHDRisk']]

gender_male=disease[disease.sex == 'male']

plt.figure(figsize=(20,20))
for i ,col in enumerate(gender_male):
```

```
plt.subplot(4,3,i+1)
ax=sns.countplot(data=gender_male,x=col ,hue='CHDRisk')
ax.bar_label(ax.containers[0])
ax.bar_label(ax.containers[1])
ax.set_title('\n\n' + col)
ax.set_xlabel('')
ax.set_ylabel('')
```



## 11 statistical tests

### 1-Chi-Square Test for Categorical Variables

- Chi-Square Test Use this to test for independence between two categorical variables. show if sex and CHDRisk are associated.

```
[408]: # Chi-square test
from scipy.stats import chi2_contingency
contingency = pd.crosstab(df['sex'], df['CHDRisk'])
chi2, p, dof, expected = chi2_contingency(contingency)
print(f"Chi-square: {chi2}, p-value: {p}")
```

```

if p < 0.05:
    print("There is a significant relationship between age_group and CHDRisk")
else:
    print("No significant relationship between age_group and CHDRisk")

```

Chi-square: 30.64922901316927, p-value: 3.091494190770845e-08

There is a significant relationship between age\_group and CHDRisk

2-Pearson Correlation Coefficient: For testing the linear relationship between two continuous variables

```


[409]: from scipy.stats import pearsonr
corr ,p_val =pearsonr(df['sysBP'],df['BMI'])
print(f'Pearson Correlation={corr}, p={p}')
if p < 0.05:
    print('There is a significant correlation between sysBP and BMI')
else:
    print('No significant correlation between sysBP and BMI')

sns.regplot(data=df,x='sysBP', y='BMI')
plt.show()

```

Pearson Correlation=0.3292817576370576, p=3.091494190770845e-08

There is a significant correlation between sysBP and BMI



## 12 Logistic Regression Significance Test:

For testing whether predictor variables in a logistic regression are significant

```
[410]: from sklearn.preprocessing import LabelEncoder
```

```
le=LabelEncoder()
df['smokingStatus']=le.fit_transform(df['smokingStatus'])
df['BPMeds']=le.fit_transform(df['BPMeds'])
df['prevalentStroke']=le.fit_transform(df['prevalentStroke'])
df['prevalentHyp']=le.fit_transform(df['prevalentHyp'])
df['diabetes']=le.fit_transform(df['diabetes'])
df['age_group'] = le.fit_transform(df['age_group'])
```

```
[413]: df['sex']=df['sex'].replace('female',0)
df['sex']=df['sex'].replace('male',1)

df['CHDRisk']=df['CHDRisk'].replace('no',0)
df['CHDRisk']=df['CHDRisk'].replace('yes',1)

df.sample(5)
```

```
[413]:
```

	sex	education	smokingStatus	cigsPerDay	BPMeds	prevalentStroke	\
2585	1	1	0	0	0	0	
425	1	3	1	20	0	0	
830	0	4	0	0	0	0	
2413	0	2	0	0	0	0	
1603	0	1	0	0	1	0	

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	\
2585	0	0	188	105.00	65.00	22.85	63	76	
425	0	0	340	134.00	89.50	21.91	50	72	
830	1	0	268	151.00	98.00	20.34	72	60	
2413	0	0	215	110.00	70.00	19.64	70	87	
1603	1	1	294	195.00	90.00	27.73	72	127	

	CHDRisk	age_group
2585	0	0
425	1	1
830	0	1
2413	0	0
1603	0	1

```
[414]: import statsmodels.api as sm

X=df.drop(columns=['CHDRisk'])
X=sm.add_constant(X)
y=df['CHDRisk']
```

```
[415]: logit_model=sm.Logit(y,X)
```

```
[416]: results=logit_model.fit()
print(results.summary())
```

Optimization terminated successfully.

Current function value: 0.381751

Iterations 7

#### Logit Regression Results

```
=====
Dep. Variable:          CHDRisk    No. Observations:          3637
Model:                Logit      Df Residuals:              3621
Method:                MLE       Df Model:                  15
Date:                  Mon, 09 Sep 2024    Pseudo R-squ.:           0.1044
Time:                  13:50:17    Log-Likelihood:          -1388.4
converged:              True      LL-Null:                 -1550.3
Covariance Type:        nonrobust    LLR p-value:             6.669e-60
=====
===
```

	coef	std err	z	P> z	[0.025
0.975]					
-----					
---					
const	-6.0047	0.636	-9.435	0.000	-7.252
-4.757					
sex	0.5872	0.108	5.416	0.000	0.375
0.800					
education	-0.0796	0.050	-1.607	0.108	-0.177
0.017					
smokingStatus	0.0172	0.157	0.109	0.913	-0.291
0.325					
cigsPerDay	0.0161	0.006	2.584	0.010	0.004
0.028					
BPMeds	0.1488	0.237	0.628	0.530	-0.316
0.613					
prevalentStroke	0.8417	0.486	1.732	0.083	-0.111
1.794					
prevalentHyp	0.2511	0.138	1.817	0.069	-0.020
0.522					
diabetes	0.1092	0.317	0.345	0.730	-0.512
0.730					
totChol	0.0030	0.001	2.720	0.007	0.001

0.005					
sysBP	0.0201	0.004	5.346	0.000	0.013
0.028					
diaBP	-0.0092	0.006	-1.434	0.152	-0.022
0.003					
BMI	0.0079	0.013	0.618	0.537	-0.017
0.033					
heartRate	-0.0048	0.004	-1.141	0.254	-0.013
0.003					
glucose	0.0073	0.002	3.243	0.001	0.003
0.012					
age_group	0.6689	0.098	6.823	0.000	0.477
0.861					

```
=====
===
```

if the p-value of a feature is less than 0.05, it means the feature is significant. . remove all features that p-value >0.05

```
[417]: X.columns
```

```
[417]: Index(['const', 'sex', 'education', 'smokingStatus', 'cigsPerDay', 'BPMeds',
        'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP',
        'diaBP', 'BMI', 'heartRate', 'glucose', 'age_group'],
        dtype='object')
```

```
[418]: X=df[['sex','cigsPerDay','totChol','glucose','sysBP','age_group']]
X=sm.add_constant(X)
y=df['CHDRisk']
logit_model=sm.Logit(y,X)
results=logit_model.fit()
print(results.summary())
```

Optimization terminated successfully.

Current function value: 0.383639

Iterations 7

#### Logit Regression Results

```
=====
Dep. Variable:          CHDRisk    No. Observations:          3637
Model:                Logit      Df Residuals:              3630
Method:                MLE       Df Model:                  6
Date:                  Mon, 09 Sep 2024    Pseudo R-squ.:          0.09996
Time:                  13:50:53    Log-Likelihood:          -1395.3
converged:              True      LL-Null:                 -1550.3
Covariance Type:        nonrobust    LLR p-value:             6.124e-64
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	-7.1055	0.408	-17.414	0.000	-7.905	-6.306



sex	0.5901	0.106	5.567	0.000	0.382	0.798
cigsPerDay	0.0154	0.004	3.756	0.000	0.007	0.023
totChol	0.0030	0.001	2.701	0.007	0.001	0.005
glucose	0.0079	0.002	4.602	0.000	0.005	0.011
sysBP	0.0205	0.002	9.669	0.000	0.016	0.025
age_group	0.7189	0.096	7.510	0.000	0.531	0.906

=====

```
[419]: results.params.index
```

```
[419]: Index(['const', 'sex', 'cigsPerDay', 'totChol', 'glucose', 'sysBP',
            'age_group'],
            dtype='object')
```

```
[420]: results.params.index[2:]
```

```
[420]: Index(['cigsPerDay', 'totChol', 'glucose', 'sysBP', 'age_group'],
            dtype='object')
```

```
[423]: for param in results.params.index[2:]:
        odds_ratio=np.exp(results.params[param])# np.exp Converts the coefficient
        ↪into the odds ratio, making the result easier to interpret.
        print(f"For each 1 unit increase in {param}, the odds of CHDRisk increase
        ↪by {round(odds_ratio, 2)} times, holding all else constant")
```

For each 1 unit increase in cigsPerDay, the odds of CHDRisk increase by 1.02 times, holding all else constant

For each 1 unit increase in totChol, the odds of CHDRisk increase by 1.0 times, holding all else constant

For each 1 unit increase in glucose, the odds of CHDRisk increase by 1.01 times, holding all else constant

For each 1 unit increase in sysBP, the odds of CHDRisk increase by 1.02 times, holding all else constant

For each 1 unit increase in age\_group, the odds of CHDRisk increase by 2.05 times, holding all else constant

## 13 Prediction

```
[435]: X=df[['sex','cigsPerDay','totChol','glucose','sysBP','age_group']]
        y=df['CHDRisk']
```

```
[436]: from sklearn.model_selection import train_test_split
        from sklearn.preprocessing import StandardScaler

        X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
        ↪2,random_state=42,stratify=y)

        scaler=StandardScaler()
```

```
x_train_scaled=scaler.fit_transform(X_train)
x_test_scaled=scaler.transform(X_test)
```

```
[456]: from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier

models={
    'logistic_Model':LogisticRegression(),
    'random_model':RandomForestClassifier(),
    'gradient_model':GradientBoostingClassifier(),
    'knn':KNeighborsClassifier(),
    'svm_model':SVC()
}
```

## 14 GridSearch : to find the best hyperparameters

```
[461]: from sklearn.model_selection import GridSearchCV

# Define the parameter grids for each model
param_grids = {
    "logistic_Model": {
        'C': [0.01, 0.1, 1, 10, 100],
        'penalty': ['l1', 'l2'],
        'solver': ['liblinear']
    },

    "random_model": {
        'n_estimators': [50, 100, 200],
        'max_depth': [None, 10, 20, 30],
        'min_samples_split': [2, 5, 10]
    },

    "gradient_model": {
        'n_estimators': [50, 100, 200],
        'learning_rate': [0.01, 0.1, 0.2],
        'max_depth': [3, 5, 7]
    },

    "knn": {
        'n_neighbors': [3, 5, 7, 9],
        'weights': ['uniform', 'distance']
    }
}
```

```

    },
    "svm_model": {
        'C': [0.1, 1, 10, 100],
        'kernel': ['linear', 'rbf'],
        'gamma': ['scale', 'auto']
    }
}

# Function to perform GridSearchCV
def perform_grid_search(model, param_grid):
    grid_search = GridSearchCV(model, param_grid, cv=5, scoring='accuracy',
    ↪n_jobs=-1)
    grid_search.fit(x_train_scaled, y_train)
    return grid_search

# GridSearchCV for each model
for name, model in models.items():
    print(f"Performing Grid Search for {name}...")
    grid_search = perform_grid_search(model, param_grids[name])
    print(f"Best parameters for {name}: {grid_search.best_params_}")
    print(f"Best accuracy for {name}: {grid_search.best_score_}")
    print("-" * 50)

```

```

Performing Grid Search for logistic_Model...
Best parameters for logistic_Model: {'C': 0.01, 'penalty': 'l2', 'solver':
'liblinear'}
Best accuracy for logistic_Model: 0.8514937511459681
-----

Performing Grid Search for random_model...
Best parameters for random_model: {'max_depth': 20, 'min_samples_split': 10,
'n_estimators': 200}
Best accuracy for random_model: 0.8483962359008939
-----

Performing Grid Search for gradient_model...
Best parameters for gradient_model: {'learning_rate': 0.01, 'max_depth': 5,
'n_estimators': 50}
Best accuracy for gradient_model: 0.8484015591083036
-----

Performing Grid Search for knn...
Best parameters for knn: {'n_neighbors': 9, 'weights': 'uniform'}
Best accuracy for knn: 0.8470258057265883
-----

Performing Grid Search for svm_model...
Best parameters for svm_model: {'C': 1, 'gamma': 'auto', 'kernel': 'rbf'}
Best accuracy for svm_model: 0.8490888443316713
-----

```

## 15 Data is Imbalanced:

```
[457]: from sklearn.metrics import confusion_matrix
model=SVC()
model.fit(x_train_scaled,y_train)
y_pred=model.predict(x_test_scaled)
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

Confusion Matrix:

```
[[617   0]
 [108   3]]
```

Total Negative Cases (TN + FP):  $617 + 0 = 617$  Total Positive Cases (FN + TP):  $108 + 3 = 111$  .

The imbalance is evident because the number of samples in the majority class (no CHD risk) is much larger than the number of samples in the minority class (CHD risk). Specifically, there are 617 instances of the majority class compared to only 111 instances of the minority class .

So ,we will use classification\_report not Confision Matrix

```
[458]: print(df['CHDRisk'].value_counts())
```

```
CHDRisk
0      3084
1       553
Name: count, dtype: int64
```

it is also here number of 0 class > number of 1 class

```
[459]: from sklearn.metrics import accuracy_score, classification_report

for name , model in models.items():
    model.fit(x_train_scaled,y_train)
    y_pred=model.predict(x_test_scaled)
    print(f"{name}:")
    print(f"Accuracy: {accuracy_score(y_test, y_pred)}")
    print(classification_report(y_test, y_pred))
    print("-" * 50)
```

logistic\_Model:

Accuracy: 0.853021978021978

	precision	recall	f1-score	support
0	0.86	0.99	0.92	617
1	0.64	0.08	0.14	111
accuracy			0.85	728
macro avg	0.75	0.54	0.53	728
weighted avg	0.82	0.85	0.80	728

-----  
random\_model:

Accuracy: 0.8475274725274725

	precision	recall	f1-score	support
0	0.86	0.97	0.92	617
1	0.50	0.15	0.23	111
accuracy			0.85	728
macro avg	0.68	0.56	0.57	728
weighted avg	0.81	0.85	0.81	728

-----  
gradient\_model:

Accuracy: 0.8502747252747253

	precision	recall	f1-score	support
0	0.86	0.99	0.92	617
1	0.56	0.08	0.14	111
accuracy			0.85	728
macro avg	0.71	0.53	0.53	728
weighted avg	0.81	0.85	0.80	728

-----  
knn:

Accuracy: 0.8543956043956044

	precision	recall	f1-score	support
0	0.87	0.98	0.92	617
1	0.57	0.18	0.27	111
accuracy			0.85	728
macro avg	0.72	0.58	0.60	728
weighted avg	0.82	0.85	0.82	728

-----  
svm\_model:

Accuracy: 0.8516483516483516

	precision	recall	f1-score	support
0	0.85	1.00	0.92	617
1	1.00	0.03	0.05	111
accuracy			0.85	728
macro avg	0.93	0.51	0.49	728
weighted avg	0.87	0.85	0.79	728

---

## 16 Summary of the Logistic Model Results:

- **Precision** tells you how accurate your positive predictions are.
- **Recall** tells you how well your model detects all actual positives.
- **F1-Score** provides a balanced measure of precision and recall, especially useful when dealing with imbalanced classes.

. - **Accuracy:** The model has an overall accuracy of **85.3%**, meaning it correctly predicts 85.3% of the instances.

- **Class 0 (majority class):**
  - **Precision:** **86%** of the instances predicted as class 0 are correct.
  - **Recall:** **99%** of the actual class 0 instances are correctly identified.
  - **F1-Score:** **92%**, indicating excellent performance for class 0.
- **Class 1 (minority class):**
  - **Precision:** **64%** of the instances predicted as class 1 are correct.
  - **Recall:** Only **8%** of the actual class 1 instances are identified correctly.
  - **F1-Score:** **14%**, showing poor performance in detecting class 1 due to the low recall.

. - **Macro Avg** (unweighted average): - **Precision:** **75%**, averaging the precision of both classes equally. - **Recall:** **54%**, reflecting the model's lower performance in detecting class 1. - **F1-Score:** **53%**, showing the balance between precision and recall.

- **Weighted Avg** (weighted by class frequency):
  - **Precision:** **82%**, with more weight given to class 0.
  - **Recall:** **85%**, influenced by the high recall of class 0.
  - **F1-Score:** **80%**, emphasizing good overall performance but issues with class 1.

### 16.0.1 Key Issue:

In your results, the model performs very well for class 0 but poorly for class 1, showing that precision and recall are not balanced for the minority class.

## 17 Compare the ROC Curves of two Models.

```
[460]: from sklearn.metrics import roc_auc_score, roc_curve

model_1=LogisticRegression()
model_2=GradientBoostingClassifier()

model_1.fit(x_train_scaled,y_train)
y_pred_model_1 = model_1.predict(x_test_scaled)
y_prob_model_1 = model_1.predict_proba(x_test_scaled)[: , 1]  # Probabilities
↳ for AUC

model_2.fit(x_train_scaled,y_train)
```

```

y_pred_model_2 = model_2.predict(x_test_scaled)
y_prob_model_2 = model_2.predict_proba(x_test_scaled)[: , 1]  # Probabilities
↳ for AUC

# Calculate ROC AUC score
roc_auc_1 = roc_auc_score(y_test, y_prob_model_1)
print(f'AUC: {roc_auc_1}')

roc_auc_2 = roc_auc_score(y_test, y_prob_model_2)
print(f'AUC: {roc_auc_2}')

# Plot ROC curve
fpr1, tpr1, thresholds1 = roc_curve(y_test, y_prob_model_1)
fpr2, tpr2, thresholds2 = roc_curve(y_test, y_prob_model_2)


plt.plot(fpr1, tpr1, color='blue', label=f'AUC = {roc_auc_1:.2f}')
plt.plot(fpr2, tpr2, color='orange', label=f'AUC = {roc_auc_2:.2f}')

plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()

```

AUC: 0.7311752595383065

AUC: 0.7025274869683298



LogisticRegression is better than gradient

What is an ROC Curve?

The ROC (Receiver Operating Characteristic) Curve is a graphical representation of the performance of a classification model at different threshold values. It plots two metrics:

True Positive Rate (Recall) on the y-axis:

This measures how many actual positives (True Positives) are correctly identified.

$$\text{True Positive Rate (TPR)} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

False Positive Rate on the x-axis:

This measures how many actual negatives are incorrectly classified as positives.

$$\text{False Positive Rate (FPR)} = \frac{\text{False Positives}}{\text{False Positives} + \text{True Negatives}}$$

Key Points:

The closer the ROC curve is to the top-left corner, the better the model's performance.

A random classifier would produce a diagonal line (FPR = TPR), representing an AUC of 0.5.



In summary, the ROC curve helps visualize the trade-off between the True Positive Rate and the False Positive Rate, while the AUC quantifies the overall performance of the model. A higher AUC means a better performing model.

Interpretation of AUC (Area Under the Curve):

AUC = 1.0: The model is perfect, distinguishing between classes with 100% accuracy.

0.9 AUC < 1.0: The model has excellent performance, meaning it is highly accurate in distinguishing between the classes.

0.8 AUC < 0.9: The model has good performance, but not perfect.

0.7 AUC < 0.8: The model has fair performance. It can distinguish between the classes better than random guessing, but there's room for improvement.

0.5 AUC < 0.7: The model has poor performance, only slightly better than random guessing.

AUC = 0.5: The model has no discriminatory power and is equivalent to random guessing.

[ ]: